


# Parallel Programming Tutorial - Pthread 1

M.Sc. Andreas Wilhelm

Technical University Munich

May 08, 2017



*TUM Uhrenturm*

# Organization

# Organization

- Tutorial starts every Monday at 4:00 PM
- Duration: as long as we need, up to 90 min
- Assignments on parallel programming techniques
- Topics
  - Pthreads (Posix Threads)
  - C++(11/14/17)
  - OpenMP (Open Multi-Processing)
  - Dependency analysis
  - MPI (Message Passing Interface)
- Code examples are in C/C++
- My email address is: `andreas.wilhelm(at)in.tum.de`

# Assignments

Starting this week, you have two weeks time for the first assignment!

- Submission of 80% of the assignments gives 0.3 bonus
- Submission server: <https://parprogr.lrr.in.tum.de/Submission>
  - requires your LRZ ID and your password
  - password is not stored and only used for authentication
- Submissions will be checked for:
  - plagiarism
  - correctness (output, threads, synchronization)
  - speedup
  - memory leaks
- Assignment instructions are on the last slides
- Final exam will contain small programming tasks (max 50% of the overall questions)
- Example solutions will be presented at the following tutorial session

# Assistance on Assignments

Starting this week

- Given by: Jeeta Chacko and Jophin John
- Email: jeetachacko(at)gmail.com / jophinjohn(at)outlook.com
- Room: 01.04.011
- Date and Time:
  - Wednesday 11:30AM - 1:00PM
  - Friday 11:30AM - 1:00PM
- If you have questions, write an email to Jeeta and Jophin or visit the assistance sessions

# Resources


- POSIX Threads Programming
- An Introduction to Parallel Programming, by Peter Pacheco
- Programming with Posix Threads, by David Butenhof
- Patterns for Parallel Programming, by Timothy G. Mattson; Beverly A. Sanders; Berna L. Massingill
- Multithreading in Modern C++, by Rainer Grimm

# Course Prerequisites

- Knowledge of C/C++
  - memory management
  - pointers /references
  - global vs. static variables
- C/C++(11/14/17) books
  - (C89) The C Programming Language, Second Edition, by Brian W. Kernighan; Dennis M. Ritchie
  - (C99) C Primer Plus, Fifth Edition, by Stephen Prata
  - (C++11/14) The C++ Programming Language, Fourth Edition, by Bjarne Stroustrup
- Experience with Linux Command Line
- Resources
  - Book: The Linux Command Line
  - Basic video introduction: The Shell
- Knowing GCC
  - An Introduction to GCC, by Brian Gough


# Year 2005: The Free Lunch Is Over

- A Fundamental Turn Toward Concurrency in Software
- Software doesn't get (much) faster with the next microprocessor generation
- Developers have to rewrite their software so that multiple computation units are used
- Parallel Programming is hard
  - to write - higher code complexity
  - to do it correctly - easy to introduce bugs
  - to debug - order of thread execution is undefined
  - to make it scalable - will your applications scale with additional cores?
- → Qualified developers are necessary



Original data collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond and C. Batten  
Dotted line extrapolations by C. Moore





# Posix Thread Programming

## Definition: Thread

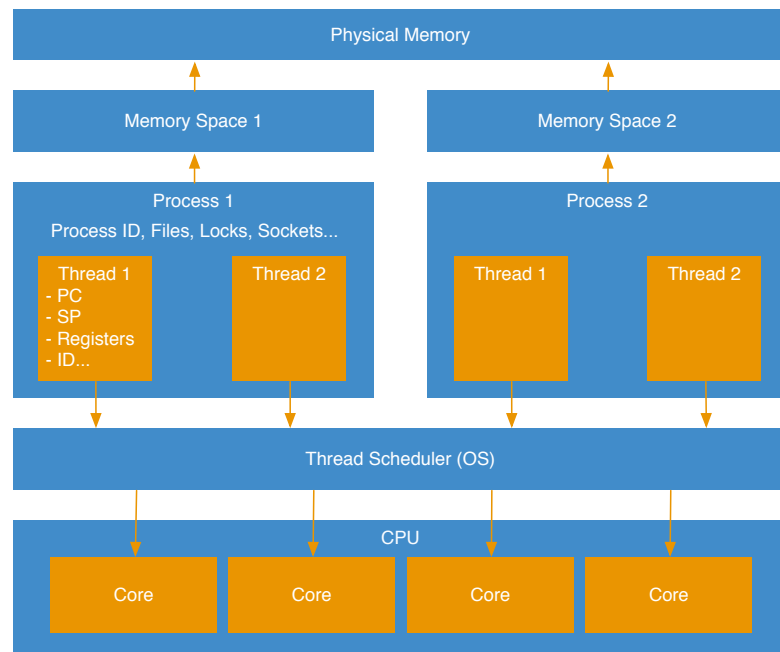
A thread is an independent stream of instructions that can be scheduled to run as such by the operating system.

## POSIX Threads (Pthreads)


- Were defined in 1995 (IEEE Std 1003.1c-1995)
- Is an API that defines a set of types, functions and constants
- Is implemented with a `pthread.h` header and a thread library
- Natively supported by FreeBSD, NetBSD, OpenBSD, Linux, Mac OS X, Android and Solaris
- Functions can be categorized in four groups:
  - Thread management
  - Mutexes
  - Condition variables *not*
  - Read/write locks and barriers

# Why use Multithreading?

- **Performance gains**  
Parallel processing by multiple processor cores
- **Increased application throughput**  
Asynchronous system calls possible
- **Increased application responsiveness**  
Application does not need to block operations
- **Replacing process-to-process communications**  
Threads may communicate by shared-memory
- **Efficient use of system resources**  
Lightweight context switches possible
- **Separation of concerns**  
Some problems are inherently concurrent



# Concurrency vs. Parallelism



## Pthread Syntax / Semantics


# Create Pthreads

```


1 int pthread_create(pthread_t *thread,
2                   const pthread_attr_t *attr,
3                   void *(*start_routine) (void *),
4                   void *arg);

```


- pthread\_t \*thread,  
– Pointer to thread identifier.
- const pthread\_attr\_t \*attr  
– Optional pointer to pthread\_attr\_t to define behavior, if NULL defaults are used.
- void \*(\*start\_routine) (void \*),  
– Pointer to function prototype that is started. Function takes void pointer as argument and returns a void pointer.
- void \*arg  
– Pointer to the argument that is used for the executed function.



# Waiting for Pthread to finish

```
1  int pthread_join(pthread_t thread,  
2 void **retval);
```

- pthread\_t thread,
  - Pointer to thread identifier, for which this function is waiting.
- void \*\*retval
  - Optional pointer pointing to a void pointer. This can be used to return data of undefined size.



# Create Pthreads - Example

```

1  #include <stdio.h>
2  #include <pthread.h>
3
4  void* hello(void* args)
5  {
6      printf("Hello World from pthread!\n");
7      return NULL;
8  }
9
10 int main()
11 {
12     pthread_t thread;
13
14     pthread_create(&thread, NULL, &hello, NULL);
15     printf("Hello World from main!\n");
16     pthread_join(thread, NULL);
17 }

```



## Compile & Output

```
gcc -pthread -Wall -o hello_world hello_world.c
```

```
Hello World from main!  
Hello World from pthread!
```

# More than One: Hello World with Pthreads Ver. 1

```
1 #include <stdlib.h>
2
3 int main()
4 {
5     int num_threads = 3;
6     pthread_t *thread = (pthread_t*) malloc(num_threads * sizeof(*thread));
7
8     for (int i = 0; i < num_threads; i++)
9         pthread_create(thread + i, NULL, &hello, NULL );
10
11     for (int i = 0; i < num_threads; i++)
12         pthread_join(thread[i], NULL );
13 }
```

# Output

```
[user]$ ./hello_world_2  
Hello World from pthread!  
Hello World from pthread!  
Hello World from pthread!
```

# Single Argument: Hello World with Pthreads Ver. 2

```
1 void * hello(void *ptr)
2 {
3     int arg = *(int*)ptr;
4     printf("Hello World from pthread %d!\n", arg);
5     return NULL;
6 }
```

# Single Argument: Hello World with Pthreads Ver. 2

```
1  int main()
2  {
3      int num_threads = 3;
4      pthread_t *thread;
5      int *arg;
6
7      thread = (pthread_t*) malloc(num_threads * sizeof(*thread));
8      arg = (int*) malloc(num_threads * sizeof(*arg));
9
10     for (int i = 0; i < num_threads; i++) {
11         arg[i] = i;
12         pthread_create(thread + i, NULL, &hello, arg + i);
13     }
14
15     for (int i = 0; i < num_threads; i++)
16         pthread_join(thread[i], NULL);
17
18     free(thread);
19     free(arg);
20 }
```

# Output

```
[user]$ ./hello_world_3  
Hello World from pthread 0!  
Hello World from pthread 1!  
Hello World from pthread 2!
```

# Many Arguments: Hello World with Pthreads Ver. 3

```
1 struct pthread_args
2 {
3     long thread_id;
4     long num_threads;
5 };
6
7 void * hello(void *ptr) {
8     struct pthread_args *arg = (struct pthread_args*)ptr;
9     printf("Hello World from %ld of %ld PID = %d TID = %li!\n",
10         arg->thread_id,
11         arg->num_threads,
12         getpid(),
13         pthread_self());
14
15     return NULL ;
16 }
```

# Many Arguments: Hello World with Pthreads Ver. 3

```
1 #include <unistd.h>
2 int main() {
3     long num_threads = 3;
4     pthread_t *thread;
5     struct pthread_args *arg;
6     thread = (pthread_t*) malloc(num_threads * sizeof(*thread));
7     arg = (struct pthread_args*) malloc(num_threads * sizeof(*arg));
8
9     for (int i = 0; i < num_threads; i++) {
10         arg[i].thread_id = i;
11         arg[i].num_threads = num_threads;
12         pthread_create(thread + i, NULL, &hello, arg + i);
13     }
14
15     for (int i = 0; i < num_threads; i++)
16         pthread_join(thread[i], NULL);
17
18     free(thread);
19     free(arg)
20 }
```



# Output

```
[user]$ ./hello_world_4  
Hello World from pthread 1 of 3 PID = 23750 TID = 23752!  
Hello World from pthread 0 of 3 PID = 23750 TID = 23751!  
Hello World from pthread 2 of 3 PID = 23750 TID = 23753!
```

# Return Result from Pthread in struct Argument

```
1 struct pthread_args
2 {
3     int in, out;
4 };
5
6 void * triple(void *ptr)
7 {
8     struct pthread_args *arg = ptr;
9     arg->out = 3 * arg->in;
10    return NULL ;
11 }
```

# Return Result from Pthread in struct Argument

```
1 int main() {
2     int num_threads = 3; pthread_t *thread;
3     struct pthread_args *arg;
4
5     thread = (pthread_t*) malloc(num_threads * sizeof(*thread));
6     arg = (struct pthread_args*) malloc(num_threads * sizeof(*arg));
7
8     for (int i = 0; i < num_threads; i++){
9         arg[i].in = i;
10        pthread_create(thread + i, NULL, &triple , arg + i);
11    }
12
13    for (int i = 0; i < num_threads; i++){
14        pthread_join(thread[i], NULL );
15        printf("Triple of %d is %d\n",
16            arg[i].in ,
17            arg[i].out);
18    }
19    free(thread);
20    free(arg);
21 }
```

# Return Result from Pthread as Pointer to Memory

```

1
2 void * triple(void *ptr) {
3
4     int *out = (int*) malloc(sizeof(*out));
5     *out = 3 * (*(int*)ptr);
6
7     return (void*)out;
8 }

```

# Return Result from Pthread as Pointer to Memory

```
1 int main() {
2     int num_threads = 3, *in;
3     pthread_t *thread;
4
5     thread = (pthread_t*) malloc(num_threads * sizeof(*thread));
6     in = (int*) malloc(num_threads * sizeof(*in));
7
8     for (int i = 0; i < num_threads; i++) {
9         in[i] = i;
10        pthread_create(thread + i, NULL, &triple, in + i);
11    }
12
13    for (int i = 0; i < num_threads; i++) {
14        int *out;
15        pthread_join(thread[i], (void*)&out);
16        printf("Triple of %d is %d\n", in[i], *out);
17        free(out);
18    }
19    free(thread);
20    free(in);
21 }
```

# What have we covered so far?


- Creating new threads with `pthread_create`
- Waiting for threads to finish with `pthread_join`
- Passing arguments to a pthread function
- Returning results from pthread function

## Assignment 1: Actors in “War and Peace”

# Assignment: histogram

Starting this week, you have two weeks time. Submission Server is not yet ready, will be announced in Moodle!

- `histogram` counts the occurrences of ten famous actors in the book “War and Peace”
- Therefore, `get_histogram` iterates over all words and compares each word with the actor names
- The task is to parallelize `get_histogram`





# Assignment: histogram

## Usage of the program

- Sequential:  
`./histogram_seq war_and_peace.txt <#threads> <#repetitions>`
- Parallel:  
`./histogram_par war_and_peace.txt <#threads> <#repetitions>`

# Assignment: histogram - get\_histogram()

```

1 void get_histogram(int nBlocks, block_t *block, int* hist, int num_threads)
2 {
3     char current_word[20] = "";
4     int c = 0;
5
6     for (int i = 0; i < nBlocks; i++) { // loop over all blocks
7         for (int j=0; j < BLOCKSIZE; j++) { // loop over all characters
8
9             if(isalpha(blocks[i][j])){ // add character to current word
10                 current_word[c++] = blocks[i][j];
11             } else { // the end of the current word
12                 current_word[c] = '\0';
13                 for(int k = 0; k < NNAMES; k++) {
14                     if(!strcmp(current_word, names[k]))
15                         hist[k]++;
16                 }
17                 c = 0;
18             }
19         }
20     }
21 }

```

# Assignment: histogram - Provided Files

- Makefile
  - contains rules to build executables
  - available targets: parallel, sequential, all (default), clean
  - 'mode=debug make [target]' to build debug version, use 'make clean' before
- main.c
  - main function - argument handling + file handling + call `get_histogram()`
- histogram.h
  - Header file for histogram.c and histogram\_\*.c
- histogram.c
  - Defines helper functions
- histogram\_seq.c
  - Sequential version of `get_histogram()`.
- student/histogram\_par.c
  - Implement the parallel version in this file

# Assignment: histogram - Provided Files

- war\_and\_peace.txt
  - Input data: The book war and peace.
- unit\_test.c
  - The unit tests that execute both the serial and parallel version to compare results.

# Assignment: Extract, Build, and Run

1. Extract all files to the current directory

```
tar -xvf assignment1.tar.gz
```

2. Build the program

```
make [sequential] [parallel] [unit_test]
```

- sequential: build the sequential program
- parallel: build the parallel program
- unit\_test: builds the unit tests

3. Run the sequential program (100 repetitions)

```
student/histogram_seq war_and_peace.txt 1 100
```

4. Run the parallel program (with N threads and 100 repetitions)

```
student/histogram_par war_and_peace.txt N 100
```

# Submission

1. Log into the website
2. Go to Assignments
3. Use the link for Assignment 1
4. Upload your `histogram_par.c` file
5. Press Submit

Parallel Programming
Home
Slides
**Assignments**
Contact
Welcome ga49qez Logout

NOTE:

1. Some of the gcc versions does not support -Wpedantic compiler flag. If you get an error saying Unrecognized command line option "-Wpedantic", change the flag to -pedantic in Makefile
2. Add -lrt option to LDFLAGS in case you get Undefined reference to "clock\_gettime" error.

Choose File
no file selected
Submit

Assignment Name	Submitted By	Date Of Submission	Status
Histogram Pthreads	Andreas Wilhelm	April 20, 2015, 6:20 p.m.	✓ Accepted

```

Build step successfull
Correctness checks successfull
Pthread checks successfull
Memcheck checks successfull
Helgrind checks successfull
Runtime checks successfull - speedup: 4.1

```